


COSC 458-647

Application Software Security

Linux Firewall

What Is A Firewall?

- A firewall is a *device* (or *software* feature) designed to control the flow of traffic into and out-of a network.
- In general, firewalls are installed to prevent attacks.



What Is A Firewall? (Cont'd)

- A mechanism to enforce **security policy**
 - Choke point that traffic has to flow through
 - Access Control List (ACLs) on a host/network level
- **Policy Decisions**
 - **What traffic should be allowed into network?**
 - **Integrity**: protect integrity of internal systems
 - **Availability**: protection from DOS attacks
 - **What traffic should be allowed out of network?**
 - **Confidentiality**: protection from data leakage

Why Use A Firewall?


- Protect a wide range of machines from general probes and many attacks.
- What is an attack?
 - Someone probing a network for computers.
 - Someone attempting to crash services on a computer.
 - Someone attempting to crash a computer.
 - Someone attempting to gain access to a computer to use resources or information.

Edge Firewall

- An edge firewall is usually software running on a server or workstation.
- An edge firewall protects a single computer from attacks directed against it.
- Examples:
 - ZoneAlarm
 - BlackIce
 - IPFW on OSX

Firewall Appliance

- An appliance firewall is a device whose sole function is to act as a firewall.
- Examples: Cisco PIX. Netscreen series.



Network Firewall

- Router/Bridge based Firewall

- A firewall running on a bridge or a router protects from a group of devices to an entire network.
- E.g., Cisco has firewall feature sets in their Internetwork Operating System (IOS) operating system.

- Computer-based Network Firewall

- A network firewall runs on a computer (such as a PC or Unix computer). These firewalls are some of the most flexible.
- Many free products are available including ipfilter, pf and netfilter and iptables (in Linux).
- Commercial products include: Checkpoint Firewall-1.
- Apple OSX includes IPFW.

How Does A Firewall Work?

- Blocks packets based on:
 - Source IP Address or range of addresses.
 - Source IP Port
 - Destination IP Address or range of addresses.
 - Destination IP Port
 - Some allow higher layers up the OSI model.
 - and other protocols.
- Common ports
 - 80 HTTP
 - 443 HTTPS
 - 20 & 21 FTP (didn't know 20 was for FTP, did you?)
 - 23 Telnet
 - 22 SSH
 - 25 SMTP

Sample Firewall Rules

- Protected server: 134.71.1.25
- Protected subnet: 134.71.1.0/24
- \$internal refers to the internal network interface on the firewall.
- \$external refers to the external network interface on the firewall.

Sample Rules

For this example, when a packet matches a rule, rule processing stops.

-- Can you spot the problem in these rules? --

1. Pass in on \$external from any proto tcp to 134.71.1.25 port = 80
2. Pass in on \$external from any proto tcp to 134.71.1.25 port = 53
3. Pass in on \$external from any proto udp to 134.71.1.25 port = 53
4. Pass in on \$external from any proto tcp to 134.71.1.25 port = 25
5. Block in log on \$external from any to 134.71.1.25
6. Block in on \$external from any to 134.71.1.0/24
7. Pass in on \$external from any proto tcp to 134.71.1.25 port = 22
8. Pass out on \$internal from 134.71.1.0/24 to any keep state

What Is A State?

- When your computer makes a connection with another computer on the network, several things are exchanged including the source and destination ports.
- In a standard firewall configuration, most inbound ports are blocked.
 - This would normally cause a problem with return traffic since the source port is randomly assigned (different from the destination port).
- A state is a **dynamic rule** created by the firewall containing the *source-destination port combination*, allowing the desired return traffic to pass the firewall.

How Many States Can A Computer Have?

- A single computer could have hundreds of states depending on the number of established connections.
- Consider a server supporting POP3, FTP, WWW and Telnet/SSH access → it could have thousands of states.
- What happens without state?
- Without state, your request for traffic would leave the firewall but the reply would be blocked.

Sample State Table.

kd2.ec.csupomona.edu - IP Filter: v3.4.28 - state top

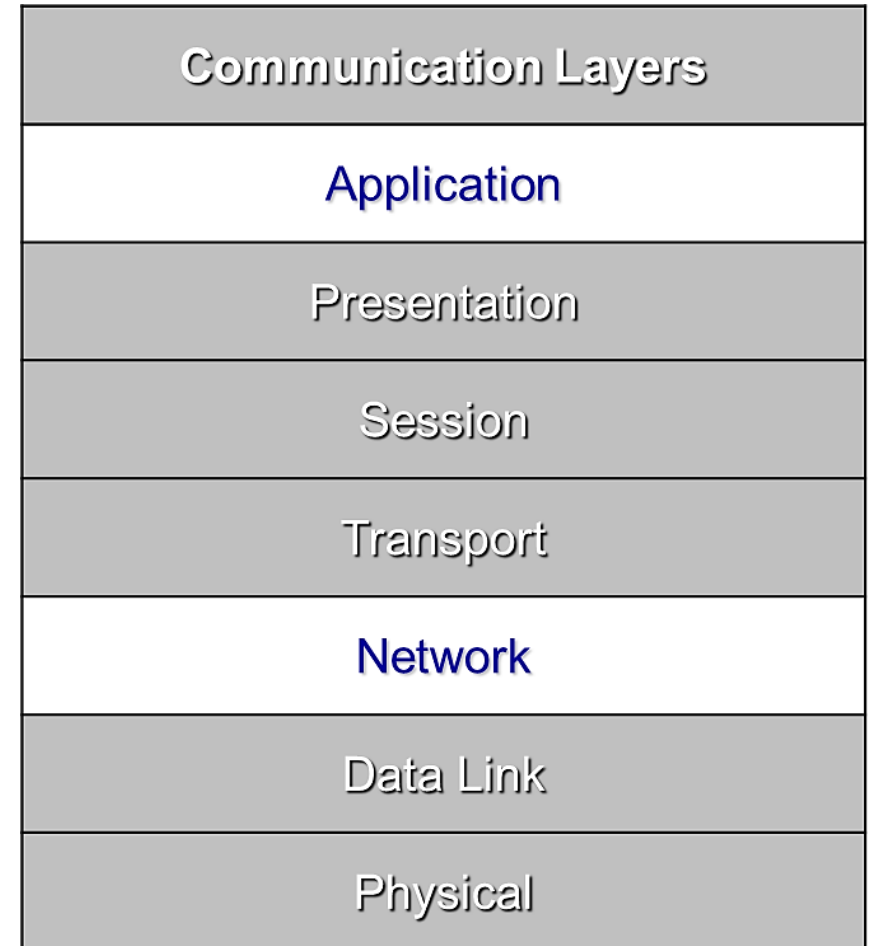
07:50:50

Src = 0.0.0.0 Dest = 0.0.0.0 Proto = any Sorted by = # bytes

Source IP	Destination IP	ST	PR	#pkts	#bytes	ttl
134.71.202.57,4738	64.160.215.222,1677	4/4	tcp	551	368024	119:59:56
134.71.202.57,4744	64.160.215.222,1677	4/4	tcp	399	258160	119:59:59
134.71.202.57,1039	134.71.204.115,1410	4/4	tcp	33	6872	119:59:16
134.71.203.168,138	134.71.203.255,138	0/0	udp	2	458	0:06
134.71.202.57,4727	64.160.215.222,1677	0/6	tcp	5	200	1:58:03
134.71.203.168,137	134.71.203.255,137	0/0	udp	2	156	0:13
134.71.202.57	239.255.255.250	0/0	igmp	1	32	1:20
134.71.202.57,137	134.71.203.255,137	0/0	udp	62	5844	1:51
134.71.202.57,1028	134.71.4.100,53	0/0	udp	35	4910	0:11
134.71.202.57,1038	216.136.175.142,5050	4/4	tcp	35	4208	119:59:59
134.71.202.57,138	134.71.203.255,138	0/0	udp	16	3520	1:49
134.71.203.168,138	134.71.203.255,138	0/0	udp	14	3026	2:00
134.71.203.168,137	134.71.203.255,137	0/0	udp	16	1536	1:59
134.71.202.57,1036	239.255.255.250,1900	0/0	udp	7	1127	1:58
134.71.202.57	239.255.255.250	0/0	igmp	10	320	1:54
134.71.202.57,4727	64.160.215.222,1677	0/6	tcp	5	200	1:53:26
134.71.202.57,1031	134.71.184.58,445	2/0	tcp	3	128	0:47
134.71.202.57,1033	134.71.184.58,445	2/0	tcp	3	128	0:48

Where Does A Firewall Fit In The Security Model?

- The firewall is the first layer of defense in any security model.
 - It should not be the only layer.
 - A firewall can stop many attacks from reaching target machines.
 - If an attack can't reach its target, the attack is defeated.
- Packet-filters work at the network layer
- Application-level gateways work at the application layer



Ruleset Design

Two main approaches to designing a ruleset are:

- Block everything then open holes.
- Block nothing then close holes.

Ruleset Design – Block Everything

- **Blocking everything** provides the strongest security but the most inconvenience.
 - Things break and people complain.
- The block everything method covers all bases ...
- ... but creates more work in figuring out how to make some applications work then opening holes.

Ruleset Design – Block Nothing

- **Blocking nothing** provides minimal security by only closing holes you can identify.
 - Blocking nothing provides the least inconvenience to our users.
- Blocking nothing means you must spend time figuring out what you want to protect yourself from then closing each hole.

Packet Filtering with netfilter

Packet Filtering

- Application-level gateways work at the application layer
- Packet-filters work at the network layer

Communication Layers
Application
Presentation
Session
Transport
Network
Data Link
Physical

Packet Filtering

- Should arriving packet be allowed in?
- Should a departing packet be let out?
- Filters packet-by-packet, forwards or drops a packet based on
 - source IP address, destination IP address
 - TCP/UDP source and destination port numbers
 - ICMP message type
 - TCP SYN and ACK bits
 - ...

Functions of Packet Filter

- Control
 - Allow only those packets that you are interested in to pass through.
- Security
 - Reject packets from malicious outsiders
- Watchfulness
 - Log packets to/from outside world

Functions of Packet Filter - Examples

- **Control** : Block incoming and outgoing datagrams with IP protocol field = 17 and with either source or dest port = 23.
- **Security**: Block inbound TCP segments with ACK=0.
 - Prevents external clients from making TCP connections with internal clients, but allows internal clients to connect to outside.

Packet Filtering in Linux

- Forward or drop packets based on TCP/IP header information, most often:
 - IP source and destination addresses
 - Protocol (ICMP, TCP, or UDP)
 - TCP/UDP source and destination ports
 - TCP Flags, especially SYN and ACK
 - ICMP message type
- Dual-homed hosts also make decisions based on:
 - Network interface the packet arrived on
 - Network interface the packet will depart on

Packet Filtering in Linux

- *netfilter* and *iptables* are the building blocks of a framework inside Linux kernel.
- *netfilter* is a set of hooks that allow kernel modules to register callback functions with the network stack.
 - Such a function is called back for every packet that traverses the respective hook.
- *iptables* is a generic table structure for the definition of rule sets.
 - Each rule within an iptable consists of a number of classifiers (iptables matches) and one connected action (iptables target).
- *netfilter*, *ip_tables*, connection tracking (*ip_conntrack*, *nf_conntrack*), and the NAT subsystem together build the whole framework.

netfilter/ iptables Capabilities

- Build Internet firewalls based on *stateless* and *stateful* packet filtering.
- Use NAT and masquerading for sharing internet access where you don't have enough addresses.
- Use NAT for implementing transparent proxies
- Mangling (packet manipulation) such as altering the TOS/DSCP/ECN bits of the IP header

netfilter

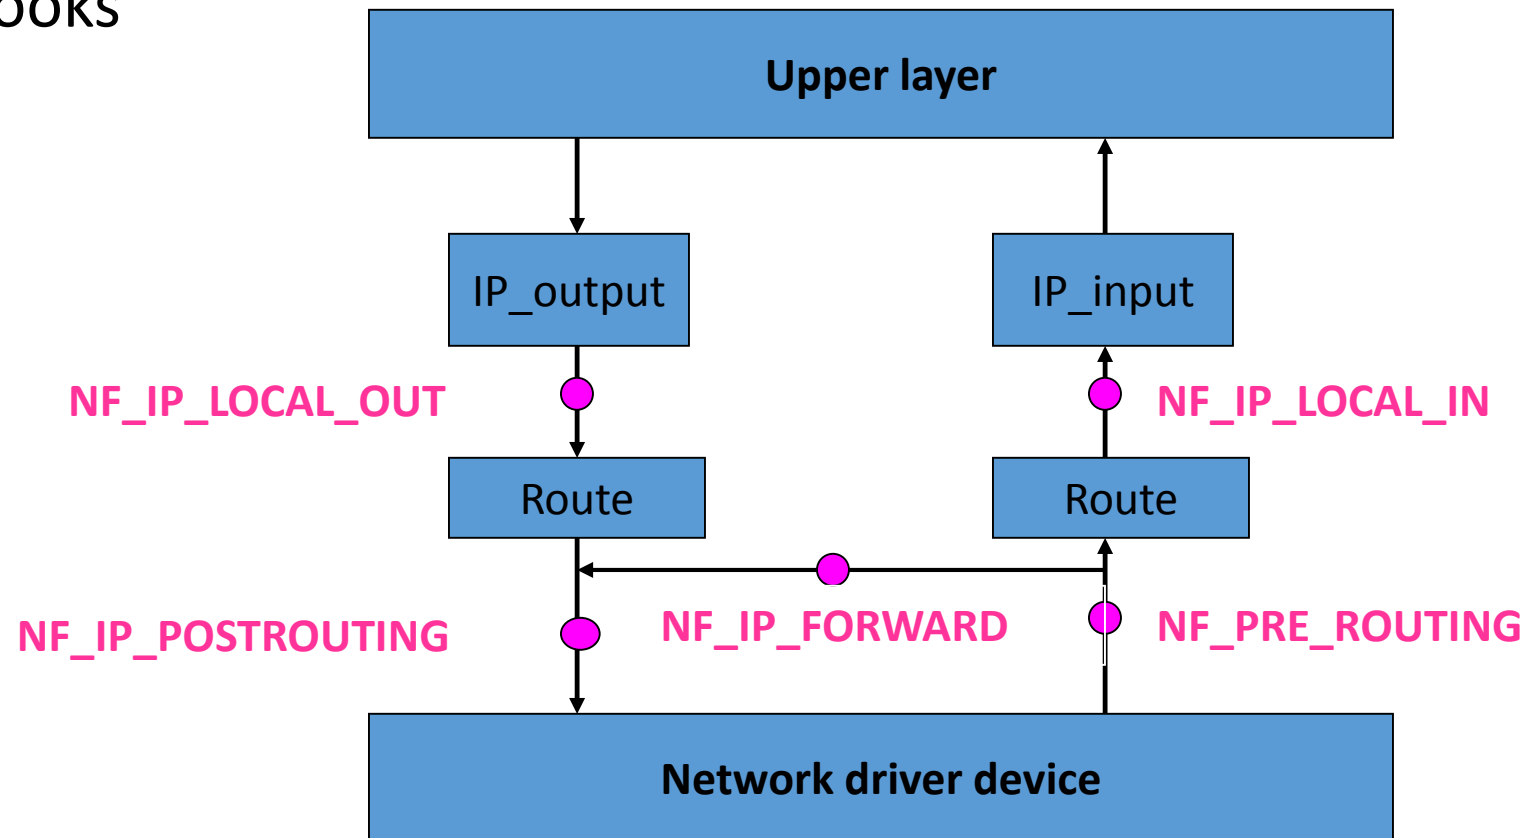
- A framework for packet mangling in kernel
- Built in Linux 2.4 kernel or higher version
- Independent of network protocol stack
- Provide an easy way to do firewall setting or packet filtering, network address translation and packet mangling

netfilter - How It Works

- Defines a set of hooks
 - Hooks are well defined point in the path of packets when these packets pass through network stack
 - The protocol code will jump into netfilter framework when it hits the hook point.
- Registers Kernel functions to these hooks
 - Called when a packet reaches at hook point
 - Can decide the fate of the packets
 - After the functions, the packet could continue its journey

Netfilter Hooks

- Each protocol family can have different hooks
 - IPv4 defines 5 hooks



Netfilter Hook Points

- **NF_IP_PRE_ROUTING:**
 - After sanity checks, before routing decision
- **NF_IP_LOCAL_IN:**
 - After routing decisions, if the packet destined for this host
- **NF_IP_FORWARD:**
 - Packets are not destined for this host but need to be forwarded to another interface
- **NF_IP_LOCAL_OUT:**
 - All packets created from local host would come here before it is sent out
- **NF_IP_POST_ROUTING:**
 - All packets have been routed and ready to hit the wire

Netfilter Hook Functions

- Multiple functions could register to the same hook point
 - Need to set priority
 - The corresponding packet will path the hook points
 - Each function registered for that hook point is called in the order of priority and is free to manipulate the packet
- What to do in hook functions
 - The function could do anything it wants
 - But need to tell netfilter to return one of the five values:
 - NF_ACCEPT
 - NF_DROP
 - NF_STOLEN
 - NF_QUEUE
 - NF_REPEAT

How To Use Netfilter Framework

- Write your own kernel functions
- Register to certain hook points
- Do whatever you want to process the packets

Add Your Own Functions To Kernel

- Two ways to extend kernel codes
- Open source programming
 - Need to find the place to insert our code
 - Recompile the whole Linux kernel
 - Need to reboot the system to let new image work
 - Time-consuming, difficult to debug
- Loadable Kernel Module (LKM)

Loadable Kernel Module

- It is a chunk of code, inserted and unloaded dynamically
- Like the normal user space programs but it works in kernel space and have access to kernel resources
- Modules take effect immediately after loading it without recompiling
- Saving time to extend the kernel

Write The Kernel Module

Different from normal C program

- Modules execute in kernel space
- May not use some standard function libraries of C
- No main functions
- Module need to define two functions with the name:
 - **init_module**: start entry
 - **cleanup_module**: end entry
 - After linux 2.4, Macro is used, we could use any name as our start and end functions, but need to use

```
module_init("start_function_name");  
module_exit("end_function_name");
```
- Use **insmod** to load, **rmmod** to unload

Example: hello.c

```
#define __KERNEL__
#define MODULE
#include <linux/module.h>
#include <linux/kernel.h>
```

```
int init_module(void) {
    printk(KERN_INFO "hello.c -- init_module() called\n");
    return 0;
}
```

```
void cleanup_module(void) {
    printk(KERN_INFO "hello.c -- cleanup_module() called\n");
}
```

```
-- Makefile --
obj-m += hello.o
LIBDIR=/lib/modules/$(shell uname -r)/build
all:
    make -C $(LIBDIR) M=$(PWD) modules
clean:
    make -C $(LIBDIR) M=$(PWD) clean
```

```
$ sudo insmod hello.ko to load into the kernel space
$ sudo rmmod hello.ko to UNload into the kernel space
$ tail /var/log/syslog to see the output
```

Netfilter Hook Implementation

1. Fill out the **nf_hook_ops** structure
2. Write a hook function:
 1. It has specific format
3. Register to system
4. Compile and load the modules

Netfilter Hook implementation

- Fill in a netfilter hook operation structure
 - Data type: (in include/linux/netfilter.h)

```
struct nf_hook_ops {  
    struct list_head list;  
    nf_hookfn *hook;           /* the callback function */  
    int pf;                   /* the network family */  
    int hooknum;              /* the hook number */  
    int priority;             /* which hook goes first */  
};
```

Example

```
static struct nf_hook_ops localin_ops;  
localin_ops.hook = <hook_func_name>;  
localin_ops.pf = PF_INET;  
localin_ops.hooknum = NF_IP_LOCAL_IN;  
localin_ops.priority = NF_IP_PRI_FIRST;
```

Netfilter Hook implementation

```
unsigned int <hook_func_name>(
    unsigned int hooknum, // the hook point where the function registered
    struct sk_buff **skb, // a reference to the packet
    const struct net_device *in_dev, // net device this packet is from/to
    const struct net_device *out_dev,
    int (*okfn) (struct sk_buff*)
)
```

Hook function returns one of the following:

NF_ACCEPT, NF_DROP, NF_STOLEN, NF_QUEUE, NF_REPEAT

Example

```
unsigned int localin_handler (
    unsigned int hook,
    struct sk_buff **skb,
    const struct net_device *indev,
    const struct net_device *outdev,
    int (*okfn) (struct sk_buff *) )
{
    struct iphdr *iphead = skb->nh.iph;
    //Drop all TCP packet
    if ( (iphead->protocol) == IPPROTO_TCP )
        return NF_DROP;
}
```

Example

- Call these functions to register/unregister at the hook
- `int nf_register_hook(struct nf_hook_ops *reg)`
- `void nf_unregister_hook(struct nf_hook_ops *reg)`

```
static int __init
init (void){
    return nf_register_hook (&localin_ops);
}

static void __exit
fini (void){
    return nf_unregister_hook (&localin_ops);
}

module_init (init);
module_exit (fini);
```

```

#include <linux/kernel.h>
#include <linux/module.h>
#include <linux/netfilter.h>
#include <linux/netfilter_ipv4.h>

static struct nf_hook_ops nfho;

//function to be called by hook
unsigned int hook_func (
    unsigned int hooknum,
    struct sk_buff **skb,
    const struct net_device *in,
    const struct net_device *out,
    int (*okfn)(struct sk_buff *))
{
    printk(KERN_INFO "packets dropped");
    return NF_DROP; //drops the packet
}

```

dropAllPackets.c

```

// Called when module loaded using 'insmod'
// Command $ (sudo insmod dropPackets.ko)
int init_module() {
    nfho.hook      = hook_func;
    nfho.hooknum   = NF_INET_PRE_ROUTING;
    nfho.pf        = PF_INET;
    nfho.priority  = NF_IP_PRI_FIRST;
    nf_register_hook(&nfho);
    printk(KERN_INFO "dropAllPacket.c --
init_module() called\n");
    return 0;
}

// Called when module unloaded using 'rmmod'
// Command $ (sudo rmmod dropPackets.ko)
void cleanup_module() {
    printk(KERN_INFO "cleanup_module()
called\n");
    //cleanup - unregister hook
    nf_unregister_hook(&nfho);
}

```

Example: dropAllPackets.c

- `dropAllPackets.c` drops all incoming packets
- Makefile:

```
obj-m += dropAllPackets.o
LIBDIR=/lib/modules/$(shell uname -r)/build
all:
    make -C $(LIBDIR) M=$(PWD) modules
clean:
    make -C $(LIBDIR) M=$(PWD) clean
```
- Load `dropAllPackets.ko` to kernel

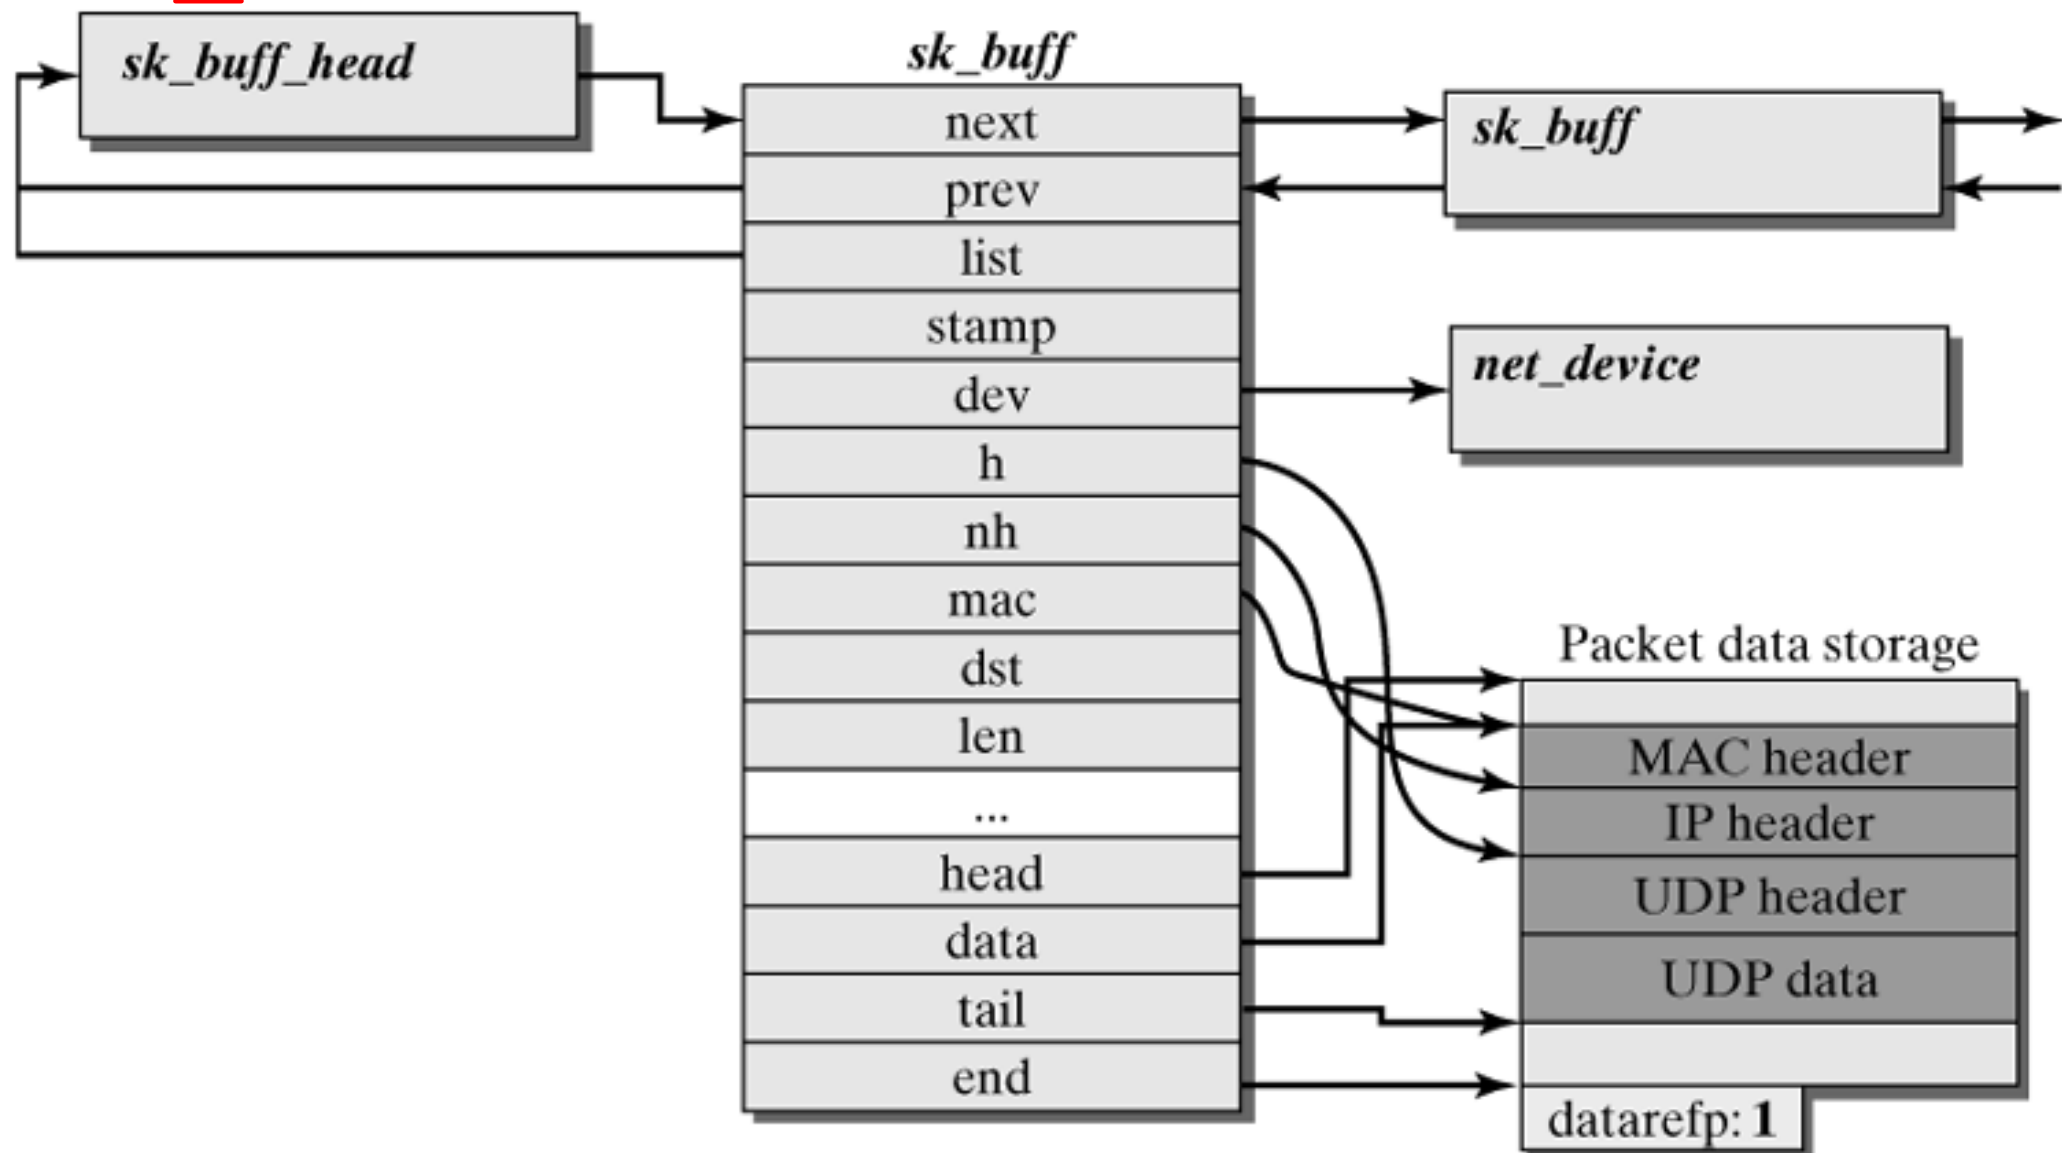
```
$ sudo insmod dropAllPackets.ko
```
- Unload `dropAllPackets.ko` from kernel

```
$ sudo rmmod hello.ko
```
- **\$ tail /var/log/syslog** to see the output

sk_buff

- Kernel buffer that stores packets.
 - Contains headers for all network layers.
- Creation
 - Application sends data to socket.
 - Packet arrives at network interface.
- Copying
 - Copied from user/kernel space.
 - Copied from kernel space to NIC.
 - Send: appends headers via `skb_reserve()`.
 - Receive: moves ptr from header to header.

sk_buff



```

struct sk_buff {
    struct sk_buff  * next;           /* Next buffer in list */
    struct sk_buff  * prev;           /* Previous buffer in list */
    struct sock *sk;                   /* Socket we are owned by */
    struct timeval  stamp;           /* Time we arrived */
    struct net_device *dev;            /* I/O net device */
    /* Transport layer header */
    union {
        struct tcphdr    *th;
        struct udphdr    *uh;
        struct icmphdr  *icmph;
        struct iphdr    *ipiph;
    } h;
    /* Network layer header */
    union {
        struct iphdr      *iph;
        struct arphdr    *arph;
    } nh;
    ...
};

```

sk_buff

Conclusion

- netfilter is a glue code between the protocol hooks and the kernel function modules
- iptables is useful to set our firewall
- Provide a simple way to hack packet

Iptables

- Based on netfilter framework
- Some modules that register to hook points
- Use generic table structure for the definition of rules.
- Very powerful
 - Customize your own firewall setting
 - User space tools (iptables) to load rules into different tables

Install iptables

- Download iptables at www.netfilter.org
- RPM
 - `rpm -ivh iptable*.i386.rpm`
- Source code
 - `tar zcvf iptable*.tar.gz`
 - `cd iptable*`
 - `./configure`
 - `make`
 - `make install`

Iptables—basic functionalities

- Packet filter
 - Control
 - Security
 - Corresponding to “filter” table
- NAT-network address translation
 - Switch the source or destination address
 - Sharing internet access
 - Corresponding to “NAT” table
- Packet mangle
 - Mangling packets going through the firewall
 - Ex: change TOS or TTL value, mark packets

Iptables command

- Use iptables command to load the rule set
- Basic iptables commands include:
 - Which table to work on
 - An operation
 - Which hook point in this table to use
 - A match
 - A target
- For example:
 - `iptables -t filter -A INPUT -p tcp -j DROP`
 - `iptables -t nat -A PREROUTING -p tcp -d 1.2.3.4 -j DNAT --to-destination 4.3.2.1`
 - `iptables -l -v -n`

To log or not to log...

Logging is both good and bad.

- If you set your rules to log too much, your logs will not be examined.
- If you log too little, you won't see things you need.
- If you don't log, you have no information on how your firewall is operating.

Sample log file

```
Jul 31 11:00:06 kd2 ipmon[14110]: 11:00:06.786765 x10 @1:10 b 134.71.4.100,50258 -> 134.71.202.57,23 PR tcp len 20 48 -S IN
Jul 31 11:00:07 kd2 ipmon[14110]: 11:00:07.366515 x10 @1:10 b 134.71.4.100,50258 -> 134.71.202.57,23 PR tcp len 20 48 -S IN
Jul 31 11:00:08 kd2 ipmon[14110]: 11:00:08.526751 x10 @1:10 b 134.71.4.100,50258 -> 134.71.202.57,23 PR tcp len 20 48 -S IN
Jul 31 11:00:10 kd2 ipmon[14110]: 11:00:10.856705 x10 @1:10 b 134.71.4.100,50258 -> 134.71.202.57,23 PR tcp len 20 48 -S IN
Jul 31 11:00:15 kd2 ipmon[14110]: 11:00:15.515785 x10 @1:10 b 134.71.4.100,50258 -> 134.71.202.57,23 PR tcp len 20 48 -S IN
Jul 31 11:50:02 kd2 ipmon[14110]: 11:50:02.619311 x10 @0:3 b 213.244.12.136,4588 -> 134.71.202.37,80 PR tcp len 20 44 -S IN
Jul 31 11:50:02 kd2 ipmon[14110]: 11:50:02.629271 x10 @0:3 b 213.244.12.136,4597 -> 134.71.202.44,80 PR tcp len 20 44 -S IN
Jul 31 11:50:02 kd2 ipmon[14110]: 11:50:02.642610 x10 @1:10 b 213.244.12.136,4610 -> 134.71.202.57,80 PR tcp len 20 44 -S IN
Jul 31 11:50:05 kd2 ipmon[14110]: 11:50:05.633338 x10 @1:10 b 213.244.12.136,4610 -> 134.71.202.57,80 PR tcp len 20 44 -S IN
Jul 31 11:50:17 kd2 ipmon[14110]: 11:50:16.882433 x10 @0:3 b 213.244.12.136,1406 -> 134.71.203.35,80 PR tcp len 20 44 -S IN
Jul 31 11:50:20 kd2 ipmon[14110]: 11:50:20.401561 x10 @0:3 b 213.244.12.136,1688 -> 134.71.203.47,80 PR tcp len 20 44 -S IN
Jul 31 11:50:20 kd2 ipmon[14110]: 11:50:20.414682 x10 @0:3 b 213.244.12.136,1701 -> 134.71.203.60,80 PR tcp len 20 44 -S IN
Jul 31 11:50:24 kd2 ipmon[14110]: 11:50:24.127364 x10 @0:3 b 213.244.12.136,1944 -> 134.71.203.103,80 PR tcp len 20 44 -S IN
Jul 31 11:50:24 kd2 ipmon[14110]: 11:50:24.144581 x10 @0:3 b 213.244.12.136,1957 -> 134.71.203.108,80 PR tcp len 20 44 -S IN
Jul 31 11:50:27 kd2 ipmon[14110]: 11:50:27.761458 x10 @0:3 b 213.244.12.136,2243 -> 134.71.203.168,80 PR tcp len 20 44 -S IN
Jul 31 11:50:27 kd2 ipmon[14110]: 11:50:27.778617 x10 @0:3 b 213.244.12.136,2260 -> 134.71.203.185,80 PR tcp len 20 44 -S IN
Jul 31 11:50:30 kd2 ipmon[14110]: 11:50:30.771581 x10 @0:3 b 213.244.12.136,2243 -> 134.71.203.168,80 PR tcp len 20 44 -S IN
Jul 31 11:50:30 kd2 ipmon[14110]: 11:50:30.772833 x10 @0:3 b 213.244.12.136,2260 -> 134.71.203.185,80 PR tcp len 20 44 -S IN
Jul 31 11:52:48 kd2 ipmon[14110]: 11:52:47.511993 x10 @1:10 b 207.45.69.69,1610 -> 134.71.202.57,113 PR tcp len 20 44 -S IN
Jul 31 11:52:51 kd2 ipmon[14110]: 11:52:50.501969 x10 @1:10 b 207.45.69.69,1610 -> 134.71.202.57,113 PR tcp len 20 44 -S IN
Jul 31 11:52:54 kd2 ipmon[14110]: 11:52:53.501498 x10 @1:10 b 207.45.69.69,1610 -> 134.71.202.57,113 PR tcp len 20 44 -S IN
Jul 31 11:52:56 kd2 ipmon[14110]: 11:52:55.703527 x10 @1:10 b 142.163.9.225,6346 -> 134.71.202.57,3343 PR tcp len 20 40 -A IN
Jul 31 11:52:57 kd2 ipmon[14110]: 11:52:56.500682 x10 @1:10 b 207.45.69.69,1610 -> 134.71.202.57,113 PR tcp len 20 44 -S IN
Jul 31 11:53:00 kd2 ipmon[14110]: 11:52:59.500694 x10 @1:10 b 207.45.69.69,1610 -> 134.71.202.57,113 PR tcp len 20 44 -S IN
Jul 31 12:00:24 kd2 ipmon[14110]: 12:00:24.220209 x10 @1:10 b 65.31.146.125,55989 -> 134.71.202.57,10336 PR tcp len 20 48 -S IN
Jul 31 12:00:26 kd2 ipmon[14110]: 12:00:26.040009 x10 @1:10 b 65.31.146.125,55989 -> 134.71.202.57,10336 PR tcp len 20 48 -S IN
Jul 31 12:00:28 kd2 ipmon[14110]: 12:00:28.794944 x10 @1:10 b 65.31.146.125,55989 -> 134.71.202.57,10336 PR tcp len 20 48 -S IN
Jul 31 12:00:34 kd2 ipmon[14110]: 12:00:34.302899 x10 @1:10 b 65.31.146.125,55989 -> 134.71.202.57,10336 PR tcp len 20 48 -S IN
Jul 31 12:00:46 kd2 ipmon[14110]: 12:00:45.284181 x10 @1:10 b 65.31.146.125,55989 -> 134.71.202.57,10336 PR tcp len 20 48 -S IN
```